

STAT 2200: Problem Set 2

Jack Ambery

Due: Monday, 2/12 at the beginning of class

- You may discuss this assignment with other students in the class, but you may not sit down and type it up with them or show them your code. You also may not discuss the assignment with anyone who isn't in our class, nor may you look up anything online.
- You must type up your homework using R Markdown. I want to see all of your code and output, and any answers you provide that aren't code must be typed above the respective R code chunk.
- Make sure none of your code runs off the page, otherwise you will lose points.
- You must print and turn in a PDF of your homework. I won't accept anything else (e.g., a Word document). All pages must be stapled together.
- This assignment is worth 100 points.

Part 1

In this part you will practice working with vectors and strings.

1. Create and store a vector that consists of the following movie lengths (in minutes): 94, 109, 110, 123, 125, 108, 92, 106, 84, 119, 110, and 140.

```
lengths <- c(94, 109, 110, 123, 125, 108, 92, 106, 84,  
119, 110, 140)
```

2. Using your entire movie length vector from problem 1, as well as the `cat()` function, reproduce the output below *exactly* as it's displayed. Note that the `round()` function allows you to round to a specified number of decimal places. Also note that you may not manually type the number 110.

Mean movie length (in min):
110 minutes

```
cat("Mean movie length (in min):\n",mean(lengths),"minutes")
```

```
## Mean movie length (in min):
```

```
## 110 minutes
```

3. Using your entire movie length vector from problem 1, as well as the `paste()` and/or `paste0()` functions, reproduce the output below *exactly* as it's displayed. You may not manually type the number 110.

“The mean movie length is 110 minutes.”

```
paste0("The mean movie length is ",mean(lengths)," minutes.")
```

```
## [1] "The mean movie length is 110 minutes."
```

Part 2

In this part you will practice working with data frames that you will create manually.

1. Create and print a data frame that consists of the following variables and values for recipes:

- type of recipe: entree, appetizer, appetizer, entree, entree, appetizer, dessert, dessert, entree, entree,
- number of ingredients: 8, 4, 5, 10, 6, 8, 7, 15, 10, 9,
- total prep time (in min): 15, 15, 5, 35, 20, 40, 25, 30, 10, 20,
- total cook time (in min): 30, 15, 20, 55, 25, 10, 120, 25, 45, 60, and
- contains meat (yes/no): yes, yes, yes, no, no, yes, no, no, yes, yes

Note: Two of the variables should be factors! Also, make sure your variable names are descriptive but not too long.

```
recipies <- data.frame(  
  type = factor(c("entree", "appetizer", "appetizer", "entree", "entree", "appetizer", "  
  ingredients = c(8, 4, 5, 10, 6, 8, 7, 15, 10, 9),  
  prepTime = c(15, 15, 5, 35, 20, 40, 25, 30, 10, 20),  
  cookTime = c(30, 15, 20, 55, 25, 10, 120, 25, 45, 60),  
  hasMeat = factor(c("yes", "yes", "yes", "no", "no", "yes", "no", "no", "yes", "yes")  
)  
)  
recipies
```

	type	ingredients	prepTime	cookTime	hasMeat
## 1	entree	8	15	30	yes
## 2	appetizer	4	15	15	yes
## 3	appetizer	5	5	20	yes
## 4	entree	10	35	55	no
## 5	entree	6	20	25	no
## 6	appetizer	8	40	10	yes
## 7	dessert	7	25	120	no
## 8	dessert	15	30	25	no
## 9	entree	10	10	45	yes
## 10	entree	9	20	60	yes

2. Create a new variable that measures the total time (prep + cook) it takes to make each recipe. Add this new variable to the data frame from problem 1, and then print the updated data frame. You may not manually type the numbers.

```
recipies$totalTime <- recipies$prepTime + recipies$cookTime  
recipies
```

	type	ingredients	prepTime	cookTime	hasMeat	totalTime
## 1	entree	8	15	30	yes	45
## 2	appetizer	4	15	15	yes	30
## 3	appetizer	5	5	20	yes	25
## 4	entree	10	35	55	no	90
## 5	entree	6	20	25	no	45
## 6	appetizer	8	40	10	yes	50
## 7	dessert	7	25	120	no	145
## 8	dessert	15	30	25	no	55
## 9	entree	10	10	45	yes	55
## 10	entree	9	20	60	yes	80

3. Create and print a second data frame that consists of the following variables and values for four additional recipes:

- type of recipe: appetizer, entree, dessert, appetizer
- number of ingredients: 3, 15, 8, 5,
- total prep time (in min): 10, 35, 45, 10,
- total cook time (in min): 0, 90, 150, 20, and
- contains meat (yes/no): no, no, no, yes

```
recipies2 <- data.frame(
  type = factor(c("appetizer", "entree", "dessert", "appetizer")),
  ingredients = c(3, 15, 8, 5),
  prepTime = c(10, 35, 45, 10),
  cookTime = c(0, 90, 150, 20),
  hasMeat = factor(c("no", "no", "no", "yes"))
)
recipies2
```

	type	ingredients	prepTime	cookTime	hasMeat
## 1	appetizer	3	10	0	no
## 2	entree	15	35	90	no
## 3	dessert	8	45	150	no
## 4	appetizer	5	10	20	yes

4. Create the same total time variable you created in problem 2 for this new data set, and then print the updated data frame.

```
recipies2$totalTime <- recipies2$prepTime + recipies2$cookTime
recipies2
```

	type	ingredients	prepTime	cookTime	hasMeat	totalTime
## 1	appetizer	3	10	0	no	10
## 2	entree	15	35	90	no	125

```
## 3  dessert      8      45      150      no      195
## 4  appetizer    5       10       20     yes       30
```

5. Combine the data frames from problems 1 and 4 so that you have a new data set with 14 observations, and then print the resulting data frame.

```
allRecipies <- rbind(recipies, recipies2)
allRecipies
```

```
##      type ingredients prepTime cookTime hasMeat totalTime
## 1  entree          8       15      30     yes         45
## 2  appetizer       4       15      15     yes         30
## 3  appetizer       5        5      20     yes         25
## 4  entree         10      35      55      no         90
## 5  entree          6      20      25      no         45
## 6  appetizer       8      40      10     yes         50
## 7  dessert        7      25     120      no        145
## 8  dessert       15      30      25      no         55
## 9  entree         10      10      45     yes         55
## 10 entree          9      20      60     yes         80
## 11 appetizer       3      10        0      no         10
## 12 entree        15      35      90      no        125
## 13 dessert        8      45     150      no        195
## 14 appetizer       5      10      20     yes         30
```

6. Print all of the data for the recipes than take under an hour to make (start to finish).

```
allRecipies[allRecipies$totalTime < 60, ]
```

```
##      type ingredients prepTime cookTime hasMeat totalTime
## 1  entree          8       15      30     yes         45
## 2  appetizer       4       15      15     yes         30
## 3  appetizer       5        5      20     yes         25
## 5  entree          6      20      25      no         45
## 6  appetizer       8      40      10     yes         50
## 8  dessert       15      30      25      no         55
## 9  entree         10      10      45     yes         55
## 11 appetizer       3      10        0      no         10
## 14 appetizer       5      10      20     yes         30
```

Part 3

In this part you will practice working with a data frame that has already been created. The data set, named `mtcars`, is included with base R.

1. What type of data structure is the data set stored as? Use an appropriate function and then type your answer above your code (outside of the R code chunk).

The `mtcars` dataset is stored as a data frame with 32 observations and 11 variables.

```
str(mtcars)

## 'data.frame':    32 obs. of  11 variables:
##  $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
##  $ cyl : num  6 6 4 6 8 6 8 4 4 6 ...
##  $ disp: num  160 160 108 258 360 ...
##  $ hp  : num  110 110 93 110 175 105 245 62 95 123 ...
##  $ drat: num  3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
##  $ wt  : num  2.62 2.88 2.32 3.21 3.44 ...
##  $ qsec: num  16.5 17 18.6 19.4 17 ...
##  $ vs  : num  0 0 1 1 0 1 0 1 1 1 ...
##  $ am  : num  1 1 1 0 0 0 0 0 0 0 ...
##  $ gear: num  4 4 4 3 3 3 3 4 4 4 ...
##  $ carb: num  4 4 1 1 2 1 4 2 2 4 ...
```

2. How many cars are in the data set? How many variables are there? Use appropriate function(s) and then type your answers above your code (outside of the R code chunk).

The `mtcars` dataset has 32 cars and 11 variables. `Columns` tells us the number of variables and `rows` tells us the number of cars.

```
# columns
ncol(mtcars)
```

```
## [1] 11
```

```
#rows
nrow(mtcars)
```

```
## [1] 32
```

3. Print the first nine rows of the data set.

```
head(mtcars, n = 9)

##           mpg  cyl  disp  hp  drat    wt  qsec vs am gear carb
## Mazda RX4      21.0    6 160.0 110  3.90  2.620 16.46  0  1    4    4
## Mazda RX4 Wag  21.0    6 160.0 110  3.90  2.875 17.02  0  1    4    4
```

```
## Datsun 710      22.8   4 108.0  93 3.85 2.320 18.61  1  1    4    1
## Hornet 4 Drive  21.4   6 258.0 110 3.08 3.215 19.44  1  0    3    1
## Hornet Sportabout 18.7   8 360.0 175 3.15 3.440 17.02  0  0    3    2
## Valiant        18.1   6 225.0 105 2.76 3.460 20.22  1  0    3    1
## Duster 360     14.3   8 360.0 245 3.21 3.570 15.84  0  0    3    4
## Merc 240D      24.4   4 146.7  62 3.69 3.190 20.00  1  0    4    2
## Merc 230       22.8   4 140.8  95 3.92 3.150 22.90  1  0    4    2
```

4. Extract the miles per gallon (mpg) variable from the data set.

```
mtcars$mpg
```

```
## [1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2 10.4
## [16] 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19.7
## [31] 15.0 21.4
```

5. Find the average and median mpg of the cars.

```
# mean
mean(mtcars$mpg)
```

```
## [1] 20.09062
```

```
# median
median(mtcars$mpg)
```

```
## [1] 19.2
```

6. Create a new data frame that consists of only the miles per gallon (mpg) and horsepower (hp) variables for all of the cars. Make sure not to remove the names of the cars.

```
mpgAndHp <- subset(mtcars, select = c(mpg, hp))
```

7. Print all of the data for the cars with at least 105 horsepower.

```
mpgAndHp[mpgAndHp$hp >= 105, ]
```

```
##           mpg  hp
## Mazda RX4    21.0 110
## Mazda RX4 Wag 21.0 110
## Hornet 4 Drive 21.4 110
## Hornet Sportabout 18.7 175
## Valiant      18.1 105
## Duster 360   14.3 245
## Merc 280     19.2 123
## Merc 280C    17.8 123
## Merc 450SE   16.4 180
```

```
## Merc 450SL          17.3 180
## Merc 450SLC         15.2 180
## Cadillac Fleetwood 10.4 205
## Lincoln Continental 10.4 215
## Chrysler Imperial   14.7 230
## Dodge Challenger    15.5 150
## AMC Javelin         15.2 150
## Camaro Z28          13.3 245
## Pontiac Firebird    19.2 175
## Lotus Europa        30.4 113
## Ford Pantera L      15.8 264
## Ferrari Dino        19.7 175
## Maserati Bora       15.0 335
## Volvo 142E          21.4 109
```

8. Print all of the data for the cars with an mpg under 20 or over 25 from the data set.

```
mpgAndHp[mpgAndHp$mpg < 20 | mpgAndHp$mpg > 25, ]
```

```
##          mpg  hp
## Hornet Sportabout 18.7 175
## Valiant           18.1 105
## Duster 360        14.3 245
## Merc 280           19.2 123
## Merc 280C         17.8 123
## Merc 450SE        16.4 180
## Merc 450SL        17.3 180
## Merc 450SLC       15.2 180
## Cadillac Fleetwood 10.4 205
## Lincoln Continental 10.4 215
## Chrysler Imperial 14.7 230
## Fiat 128           32.4  66
## Honda Civic       30.4  52
## Toyota Corolla    33.9  65
## Dodge Challenger   15.5 150
## AMC Javelin       15.2 150
## Camaro Z28        13.3 245
## Pontiac Firebird   19.2 175
## Fiat X1-9          27.3  66
## Porsche 914-2      26.0  91
## Lotus Europa      30.4 113
## Ford Pantera L     15.8 264
## Ferrari Dino       19.7 175
## Maserati Bora      15.0 335
```

9. Print all of the data for the cars with an mpg of at least 22 and a horsepower under 95.


```
mpgAndHp[mpgAndHp$mpg >= 22 & mpgAndHp$hp < 95, ]
```

```
##           mpg hp
## Datsun 710  22.8 93
## Merc 240D   24.4 62
## Fiat 128    32.4 66
## Honda Civic 30.4 52
## Toyota Corolla 33.9 65
## Fiat X1-9   27.3 66
## Porsche 914-2 26.0 91
```